

In [8]: *# Importing dependencies*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

import warnings
warnings.filterwarnings('ignore')

data = pd.read_csv("heart.csv")

data.head(10)
```

Out[8]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0
5	58	0	0	100	248	0	0	122	0	1.0	1	0	2	1
6	58	1	0	114	318	0	2	140	0	4.4	0	3	1	0
7	55	1	0	160	289	0	0	145	1	0.8	1	1	3	0
8	46	1	0	120	249	0	0	144	0	0.8	2	0	3	0
9	54	1	0	122	286	0	0	116	1	3.2	1	2	2	0

In [9]: data.describe()

Out[9]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpe
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.529756	149.114146	0.336585	1.0715
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.527878	23.005724	0.472772	1.1750
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.0000
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	132.000000	0.000000	0.0000
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	152.000000	0.000000	0.8000
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.000000	166.000000	1.000000	1.8000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.2000



```
In [10]: data.shape, data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```


```
Out[10]: ((1025, 14), None)
```

```
In [11]: data['target'].value_counts()
```


```
Out[11]: target
1      526
0      499
Name: count, dtype: int64
```

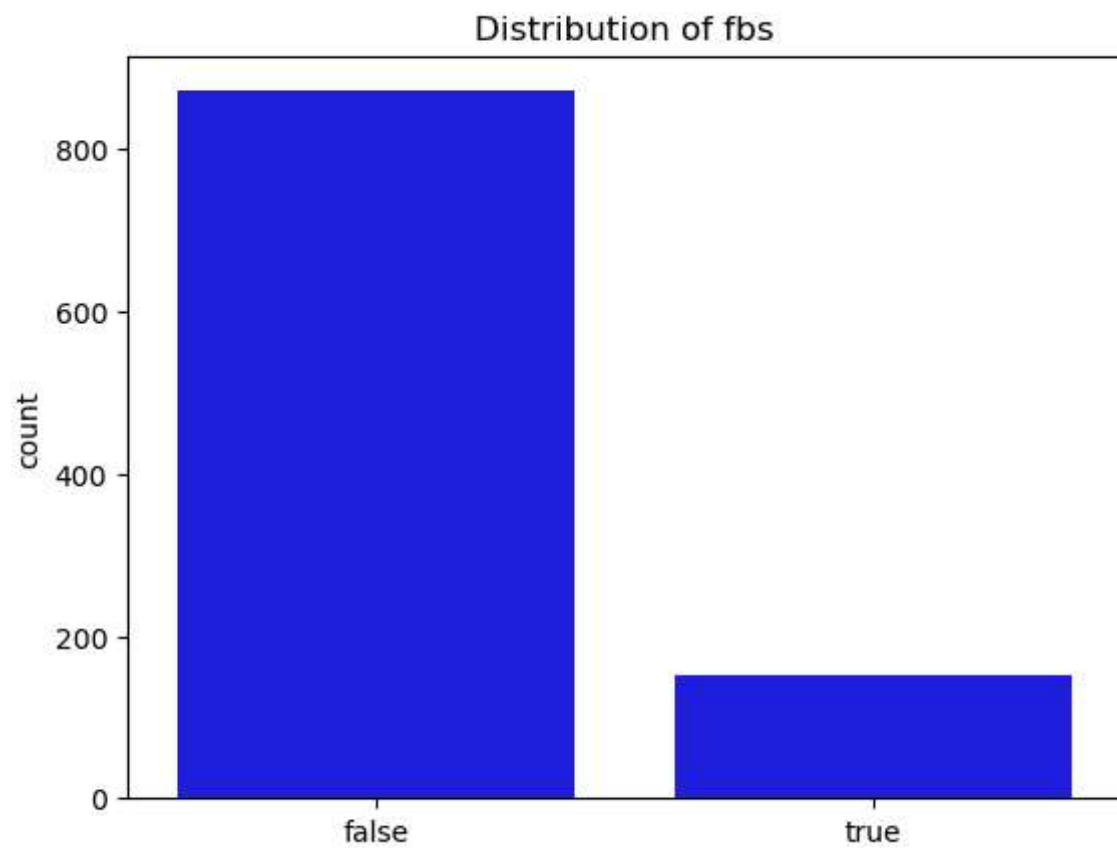
```
In [14]: # More patient values w/ heart disease than there are without.  
# Visualizing values.
```

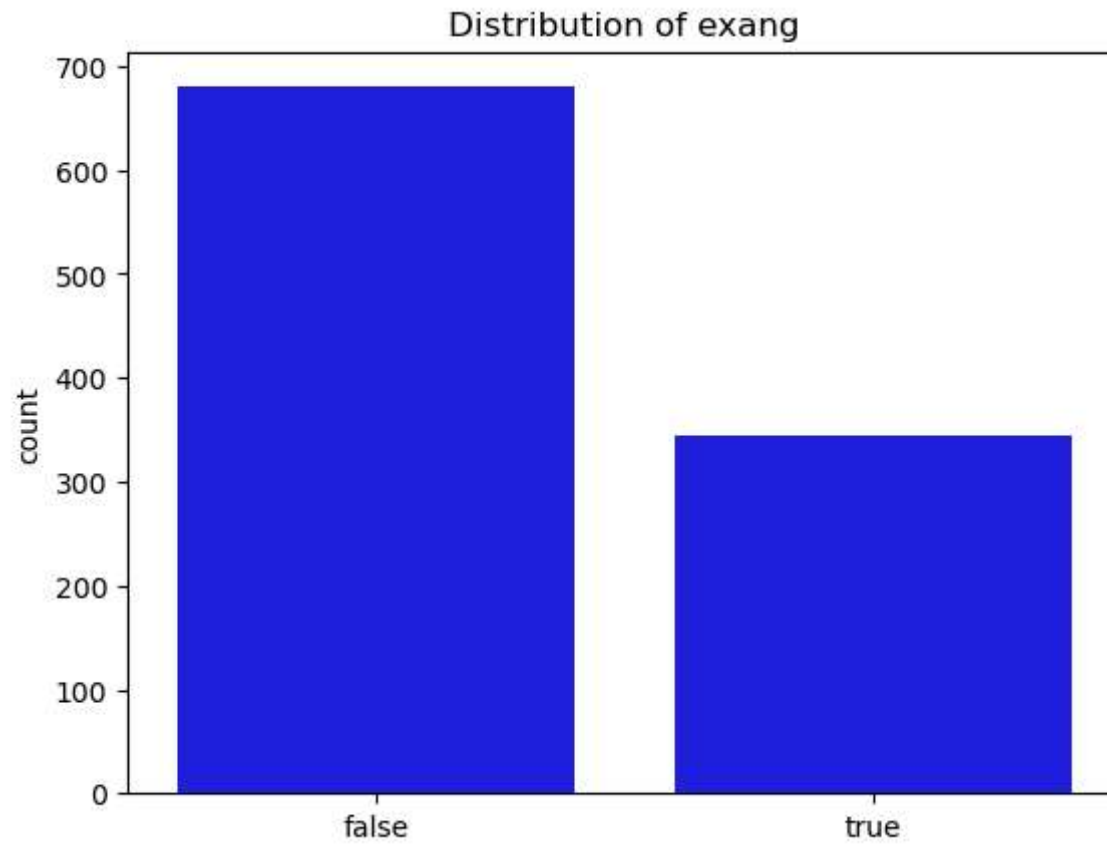
```
x = ['1', '0']  
y = data['target'].value_counts()  
  
sns.barplot(x=x, y=y, color='blue')  
plt.title('Distribution of Target Variable')  
plt.show()
```



```
In [15]: def showBinaryFeatureValues(feature1, feature2, feature):  
    x = [feature1, feature2]  
    y = data[feature].value_counts()  
  
    sns.barplot(x=x, y=y, color='blue')  
    plt.title(f'Distribution of {feature}')  
    plt.show()  
  
showBinaryFeatureValues('male', 'female', 'sex')  
showBinaryFeatureValues('false', 'true', 'fbs')  
showBinaryFeatureValues('false', 'true', 'exang')
```







```
In [58]: # Data Preprocessing
# All of the values are already numerical so there is no need for any encoding.

X = data.drop(['target'], axis=1)
Y = data['target']

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)
```

```
In [ ]: # RandomForest Classification
rfc = RandomForestClassifier(n_estimators=10, max_depth=None, min_samples_split=2,
                             max_leaf_nodes=None, min_samples_leaf=1, max_samples=None, max_features=None)

rfc.fit(x_train, y_train)

predictions = rfc.predict(x_test)
```

```
In [62]: print("Accuracy:", metrics.accuracy_score(y_test, predictions))
```

Accuracy: 0.975609756097561

```
In [63]: # To make the search for the optimal # of trees, we can iterate through a list of range of values and return

def ListComp():
    return [x for x in range(1, 100)]

n_estimators = ListComp()

accuracies = []

for i in n_estimators:
    rfc = RandomForestClassifier(n_estimators=i, max_depth=None, min_samples_split=2, max_leaf_nodes=None, min_samples_leaf=1, max_samples=None, max_features=None)
    rfc.fit(x_train, y_train)
    predictions = rfc.predict(x_test)
    accuracies.append(metrics.accuracy_score(y_test, predictions))

#shows the maximum accuracy
max_acc = max(accuracies)
print(max_acc)

# Using sklearn's classification report to see the precision, recall, and f1 score for our model.
classification_report(y_test, predictions)
```

1.0

```
Out[63]:
```

		precision	recall	f1-score	support		0	1.00	0.97	0.98
89		1	0.97	1.00	0.99	116				
205	macro avg	0.99	0.98	0.99	205	accuracy				0.99
205	weighted avg					weighted avg	0.99	0.99	0.99	

In []: