

```
In [2]: #Import python packages
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split # Import train_test_split function
from sklearn import svm #Import svm model
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy score
```

```
In [3]: #Import the heart data
data = pd.read_csv("heart.csv")
```

```
In [4]: #Display first 5 lines of heart data
data.head()
```

Out[4]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

```
In [5]: #Display basic info about the data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [6]: #Separate Feature and Target Matrix
x = data.drop('target',axis = 1)
y = data.target
```

```
In [7]: # Split dataset into training set and test set
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,random_state=109) # 70% training and
```

```
In [8]: # Create a pipeline with feature scaling and SVM classifier
pipeline = make_pipeline(StandardScaler(), svm.SVC(kernel='linear'))

# Define the parameter grid for hyperparameter tuning
param_grid = {'svc__C': [0.1, 1, 5, 10, 50], 'svc__gamma': [1, 0.1, 0.01, 0.001]}

# Perform GridSearchCV for hyperparameter tuning
grid_search = GridSearchCV(pipeline, param_grid, cv=5)
grid_search.fit(x_train, y_train)

# Get the best model from GridSearchCV
best_model = grid_search.best_estimator_

# Train the best model on the entire training set
best_model.fit(x_train, y_train)

# Predict the response for the test dataset
y_pred = best_model.predict(x_test)

# Calculate the accuracy of the best model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.8636363636363636

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.81	0.85	150
1	0.83	0.92	0.87	158
accuracy			0.86	308
macro avg	0.87	0.86	0.86	308
weighted avg	0.87	0.86	0.86	308

```
In [20]: #Create a svm Classifier  
ml = svm.SVC(kernel='linear') # Linear Kernel  
  
#Train the model using the training sets  
ml.fit(x_train, y_train)  
  
#Predict the response for test dataset  
y_pred = ml.predict(x_test)
```

```
In [21]: # Model Accuracy: how often is the classifier correct?  
ml.score(x_test, y_test)
```

Out[21]: 0.8733766233766234

In [ ]:

In [ ]: